Package 'SPARRAfairness'

April 9, 2025

Title Analysis of Differential Behaviour of SPARRA Score Across Demographic Groups

Version 0.1.0.0

Maintainer James Liley <james.liley@durham.ac.uk>

Description The SPARRA risk score (Scottish Patients At Risk of admission and Re-Admission) estimates yearly risk of emergency hospital admission using electronic health records on a monthly basis for most of the Scottish population. This package implements a suite of functions used to analyse the behaviour and performance of the score, focusing particularly on differential performance over demographically-defined groups. It includes useful utility functions to plot receiver-operator-characteristic, precision-recall and calibration curves, draw stock human figures, estimate counterfactual quantities without the need to re-compute risk scores, to simulate a semi-realistic dataset. Our manuscript can be found at Thoma, I et al: <doi.org/10.1371/journal.pdig.0000675>.

License GPL (>=3)

Encoding UTF-8

LazyData true

Depends R (>= 3.5.0), stats, graphics, grDevices, matrixStats, ranger

Imports mvtnorm,cvAUC,ggplot2,ggrepel,patchwork,scales

Roxygen list(markdown = TRUE)

RoxygenNote 7.3.2

Suggests knitr, rmarkdown

VignetteBuilder knitr

NeedsCompilation no

Author Ioanna Thoma [aut] (<https://orcid.org/0000-0001-6928-2198>),
Catalina Vallejos [ctb] (<https://orcid.org/0000-0003-3638-1960>),
Louis Aslett [ctb] (<https://orcid.org/0000-0003-2211-233X>),
Jill Ireland [ctb] (<https://orcid.org/0009-0009-5324-6630>),
Simon Rogers [ctb] (<https://orcid.org/0000-0003-3578-4477>),
James Liley [cre, aut] (<https://orcid.org/0000-0002-0049-8238>)

Contents

1

---

ab   *ab() Shorthand to draw a red x-y line*

Description

ab() Shorthand to draw a red x-y line

Usage

```
ab(...)
```

Arguments

...   passed to abline()

Value

No return value, draws a figure

---

adjusted_fdr *adjusted_fdr*

Description

Estimates false discovery rate P(target=FALSE|score>cutoff,group=g) 'adjusted' for some category.

Usage

```
adjusted_fdr(
  scores,
  target,
  category,
  group1,
  group2,
  cutoffs = seq(min(scores, na.rm = TRUE), max(scores, na.rm = TRUE), length = 100),
  nboot = 100
)
```

Arguments

scores	vector of risk scores
target	vector of values of target (which risk score aims to predict)
category	vector of categories
group1	indices of group 1
group2	indices of group 2
cutoffs	score cutoffs at which to estimate metric (default 100 evenly-spaced)
nboot	number of bootstrap samples for standard error

Details

Namely, calculates

sum (P(target=FALSE|score>cutoff,category=c,group=g)P(category=c|score<cutoff))

where the sum is over categories c.

Value

matrix of dimension length(cutoffs)x4, with (i,2g-1)th entry the relevant fairness metric for group g at the ith cutoff value and (i,2g)th entry the approximate standard error of the (i,2g-1)th entry

Examples

```
# See vignette
```

adjusted_for	*adjusted_for*

Description

Estimates false omission rate P(target=TRUE|score<=cutoff,group=g) 'adjusted' for some category.

Usage

```
adjusted_for(
  scores,
  target,
  category,
  group1,
  group2,
  cutoffs = seq(min(scores, na.rm = TRUE), max(scores, na.rm = TRUE), length = 100),
  nboot = 100
)
```

Arguments

scores	vector of risk scores
target	vector of values of target (which risk score aims to predict)
category	vector of categories
group1	indices of group 1
group2	indices of group 2
cutoffs	score cutoffs at which to estimate metric (default 100 evenly-spaced)
nboot	number of bootstrap samples for standard error

Details

Namely, calculates

sum (P(target=TRUE|score<=cutoff,category=c,group=g)P(category=c|score<cutoff))

where the sum is over categories c.

Value

matrix of dimension length(cutoffs)x4, with (i,2g-1)th entry the relevant fairness metric for group g at the ith cutoff value and (i,2g)th entry the approximate standard error of the (i,2g-1)th entry

Examples

```
# See vignette
```

---

all_data	*All data for fairness measures*

Description

This object contains all data from analysis of fairness measures in SPARRA v3 and v4.

Usage

```
all_data
```

Format

An object of class list of length 1261.

---

build_diff	*build_diff Prepares a data frame for a ggplot object to compare differences using linear interpolation.*

Description

build_diff Prepares a data frame for a ggplot object to compare differences using linear interpolation.

Usage

```
build_diff(df, xvar)
```

Arguments

df	data frame
xvar	name of variable to consider as 'x': interpolate over evenly spaced values of this variable.

Value

data frame using (common) interpolated x values rather than arbitrary x values

Examples

```
# Only used internally
```

---

cal_2panel   *cal_2panel Draws calibration curves (with legend) with a second panel underneath showing predicted differences.*

Description

cal_2panel Draws calibration curves (with legend) with a second panel underneath showing predicted differences.

Usage

```
cal_2panel(
  cals,
  labels,
  col = 1:length(cals),
  xy_col = phs_colours("phs-magenta"),
  ci_col = col,
  highlight = NULL,
  yrange_lower = NULL,
  legend_title = ""
)
```

Arguments

cals	list of calibration objects, output from getcal().
labels	labels to use in legend
col	line colours
xy_col	line colour for x-y line, defaults to phs-magenta
ci_col	colours to draw confidence intervals on lower panel; NA to not draw.
highlight	if non-null, highlight a particular value
yrange_lower	y range for lower plot. If NULL, generates automatically
legend_title	title for legend, defaults to nothing

Value

Silently return ggplot object

Examples

```
# See vignette
```

---

counterfactual_yhat   *counterfactual_yhat*

Description

Estimation of counterfactual quantities by resampling.

Usage

```
counterfactual_yhat(dat, X, x = NULL, G, g, gdash, excl = NULL, n = NULL)
```

Arguments

dat	data frame containing variables in X, Yhat, G, excl. Variables in U are assumed to be colnames(dat)\(X U Yhat U G U excl)
X	set of variables and values on which to 'condition'; essentially we assume that any causal pathway from G to Yhat is through X
x	values of variables X on which to condition; can be null in which case we use marginal distribution of X
G	grouping variable, usually a sensitive attribute
g	conditioned value of g
gdash	counterfactual value of g
excl	variable names to exclude from U
n	number of samples; if NULL return all

Details

Counterfactual fairness is with respect to the causal graph:

G <————U | / | | / | | / | VV V X ——> Yhat

where

- G=group (usually sensitive attribute);
- Yhat=outcome;
- X=set of variables through which G can act on Yhat,
- U=set of background variables;

We want the counterfactual $Yhat_{g'<-G}|X = x, G = g$ (or alternatively $(Yhat_{g'<-G}|G = g)$), using the term $Yhat$ as it appears in the function, rather than $\hat{Y}$.

This can be interpreted as: the distribution of values of Yhat amongst indivdiduals whose values of U are

Essentially, comparison of the counterfactual quantity above to the conditional $(Yhat|G = g)$ isolates the difference in Yhat due to the effect of G on Yhat through X, removing any effect due to different distributions of U due to different values of G.

To estimate $Y' = Yhat_{g' \to G}|G = g$, we need to

1. Compute $U' \sim (U|G = g)$
2. Compute the distribution X' as $X' \sim (X|U\ U', G = g')$
3. Sample Y' $(Yhat|X \sim X', U \sim U')$

To estimate $Y' = Yhat_{g' \to G}|X = x, G = g$, we need to

1. Compute $U' \sim (U|G = g, X = x)$
2. Compute the distribution X' as $X' \sim (X|U \sim U', G = g')$
3. Sample $Y' \sim (Yhat|X \sim X', U \sim U')$

This function approximates this samplying procedure as follows

1. Look at individuals with $G = g$ (and optionally $X = x$)
2. Find the values of U for these individuals
3. Find a second set of individuals with the same values of U but for whom $G = g'$
4. Return the indices of these individuals

The values of Yhat for these individuals constitute a sample from the desired counterfactual.

Value

indices representing sample(s) from counterfactual Yhat(g' <- G) |X=x,G=g

Examples

```
set.seed(23173)
N=10000

# Background variables sampler
background_U=function(n) runif(n) # U~U(0,1)

# Structural equations
struct_G=function(u,n) rbinom(n,1,prob=u) # G|U=u ~ Bern(u)
struct_X=function(u,g,n) rbinom(n,1,prob=u*(0.5 + 0.5*g)) # X|U=u,G=g ~ Bern(u(1+g)/2)
struct_Yhat=function(u,x,n) (runif(n,0,x) + runif(n,0,u))/2 # Yhat|X,N ~ (U(0,X) + U(0,U))/2

# To see that the counterfactual 'isolates' the difference in Yhat due to the
#  causal pathway from G to Yhat through X, change the definition of struct_G to
#
#  struct_G=function(u,n) rbinom(n,1,prob=1/2) # G|U=u ~ Bern(1/2)
#
# so the posterior of U|G=g does not depend on g. Note that, with this definition, the
#  counterfactual Yhat{G<-1}|G=1 coincides with the conditional Yhat|G=0, since
#  the counterfactual G<-1 is equivalent to just conditioning on G=1.
#
# By contrast, if we change struct_G back to its original definition, but
#  change the definition of struct_Yhat to
#
# struct_Yhat=function(u,x,n) (runif(n,0,1) + runif(n,0,u))/2 # Yhat|X,N ~ (U(0,1) + U(0,U))/2
#
# so Yhat depends on G only through the change in posterior of U from changing g,
#  the counterfactual Yhat{G<01}|G=1 coincides with the conditional Yhat|G=1.


# Sample from complete causal model
U=background_U(N)
G=struct_G(U,N)
X=struct_X(U,G,N)
Yhat=struct_Yhat(U,X,N)
dat=data.frame(U,G,X,Yhat)
```

```
# True counterfactual Yhat{G <- 0}|G=1
w1=which(dat$G==1)
n1=length(w1)
UG1=dat$U[w1] # This is U|G=1
XG1=struct_X(UG1,rep(0,n1),n1)
YhatG1=struct_Yhat(UG1,XG1,n1)

# Estimated counterfactual Yhat{G <- 0}|G=1
ind_G1=counterfactual_yhat(dat,X="X",G="G",g = 1, gdash = 0)
YhatG1_resample=dat$Yhat[ind_G1]




# True counterfactual Yhat{G <- 0}|G=1,X=1
w11=which(dat$G==1 & dat$X==1)
n11=length(w11)
UG1X1=dat$U[w11] # This is U|G=1,X=1
XG1X1=struct_X(UG1X1,rep(0,n11),n11)
YhatG1X1=struct_Yhat(UG1X1,XG1X1,n11)

# Estimated counterfactual Yhat{G <- 0}|G=1
ind_G1X1=counterfactual_yhat(dat,X="X",G="G",g = 1, gdash = 0,x=1)
YhatG1X1_resample=dat$Yhat[ind_G1X1]




# Compare CDFs
x=seq(0,1,length=1000)
oldpar = par(mfrow=c(1,2))

plot(0,type="n",xlim=c(0,1),ylim=c(0,1),xlab="Value",
     ylab=expression(paste("Prop. ",hat('Y')," < x")))
lines(x,ecdf(dat$Yhat)(x),col="black") # Unconditional CDF of Yhat
lines(x,ecdf(dat$Yhat[which(dat$G==1)])(x),col="red") # Yhat|G=1
lines(x,ecdf(dat$Yhat[which(dat$G==0)])(x),col="blue") # Yhat|G=0

# True counterfactual Yhat{G <- 0}|G=1
lines(x,ecdf(YhatG1)(x),col="blue",lty=2)

# Estimated counterfactual Yhat{G <- 0}|G=1
lines(x,ecdf(YhatG1_resample)(x),col="blue",lty=3)

legend("bottomright",
       c(expression(paste(hat('Y'))),
         expression(paste(hat('Y'),"|G=1")),
         expression(paste(hat('Y'),"|G=0")),
         expression(paste(hat(Y)[G %<-% 0],"|G=1 (true)")),
         expression(paste(hat(Y)[G %<-% 0],"|G=1 (est.)"))),
       col=c("black","red","blue","blue","blue"),
       lty=c(1,1,1,2,3),
       cex=0.5)
```

```
plot(0,type="n",xlim=c(0,1),ylim=c(0,1),xlab="Value",
      ylab=expression(paste("Prop. ",hat('Y'),") < x")))
lines(x,ecdf(dat$Yhat[which(dat$X==1)])(x),col="black") # CDF of Yhat|X=1
lines(x,ecdf(dat$Yhat[which(dat$G==1 & dat$X==1)])(x),col="red") # Yhat|G=1,X=1
lines(x,ecdf(dat$Yhat[which(dat$G==0 & dat$X==1)])(x),col="blue") # Yhat|G=0,X=1

# True counterfactual Yhat{G <- 0}|G=1,X=1
lines(x,ecdf(YhatG1X1)(x),col="blue",lty=2)

# Estimated counterfactual Yhat{G <- 0}|G=1,X=1
lines(x,ecdf(YhatG1X1_resample)(x),col="blue",lty=3)

legend("bottomright",
       c(expression(paste(hat('Y|X=1'))),
         expression(paste(hat('Y'),"|G=1,X=1")),
         expression(paste(hat('Y'),"|G=0,X=1")),
         expression(paste(hat(Y)[G %<-% 0],"|G=1,X=1 (true)")),
         expression(paste(hat(Y)[G %<-% 0],"|G=1,X=1 (est.)"))),
       col=c("black","red","blue","blue","blue"),
       lty=c(1,1,1,2,3),
       cex=0.5)

# In both plots, the estimated counterfactual CDF closely matches the CDF of the
#  true counterfactual.

# Restore parameters
par(oldpar)
```

---

dat2mat   *dat2mat*

Description

Generates matrices for decomposition of admission type which can be used in `plot_decomp`

Usage

```
dat2mat(dat, score, group1, group2, nquant = 20, cats = unique(dat$reason))
```

Arguments

dat	data frame with population data, such as output from sim_pop_data. Must include a column `reason`
score	risk scores corresponding to dat
group1	indices for group 1
group2	indices for group 2
nquant	number of quantiles of code to use; default 20
cats	vector of strings giving names of admission categories; default the unique values in dat$reason. Can include NAs.

Details

Generates two matrices with the following specifications: Each matrix corresponds to one group Columns are named with the admission types to be plotted. Any admission types including the string 'Died' are counted as deaths If the matrix has N rows, these are interpreted as corresponding to N score quantiles The (i,j)th entry of the matrix is the number of people admitted for reason i with a score greater than or equal to (j-1)/N and less than (j/N) who are in that group

Value

list with two objects `matrix1` and `matrix2` giving output matrices

Examples

```
# See vignette
```

---

decomposition_matrix *Decomposition matrix*

Description

Matrix giving frequency of admission types for various groups at various score thresholds. Row names are of the form vX_Y_qZ, where X is version (3 or 4), Y is cohort (e.g., all, over 65, island postcode) and Z is quantile (1-20) of score. Column names are cause of admission or cause of death.

Usage

```
decomposition_matrix
```

Format

An object of class `data.frame` with 520 rows and 41 columns.

---

demographic_parity *demographic_parity*

Description

Estimates demographic parity for a risk score (essentially cumulative distribution function)

Usage

```
demographic_parity(
  scores,
  group1,
  group2,
 cutoffs = seq(min(scores, na.rm = TRUE), max(scores, na.rm = TRUE), length = 100)
)
```

Arguments

scores	vector of risk scores
group1	indices of group 1
group2	indices of group 2
cutoffs	score cutoffs at which to estimate DP (default 100 evenly-spaced)

Value

matrix of dimension length(cutoffs)x4, with (i,2g-1)th entry the proportion of scores in group g which are less than or equal to the ith cutoff value and (i,2g)th entry the approximate standard error of the (i,2g-1)th entry

Examples

```
# See vignette
```

---

drawperson   *drawperson*

Description

Draws a simple stock image of a person.

Usage

```
drawperson(
  xloc = 0,
  yloc = 0,
  scale = 1,
  headsize = 0.16,
  headangle = pi/8,
  headloc = 0.5,
  necklength = 0.1,
  shoulderwidth = 0.1,
  shouldersize = 0.05,
  armlength = 0.4,
  armangle = 7 * pi/8,
  armwidth = 0.08,
  leglength = 0.5,
  legangle = 9 * pi/10,
  legwidth = 0.15,
  torsolength = 0.4,
  ...
)
```

Arguments

xloc	x-axis offset from origin
yloc	y-axis offset from origin
scale	scale upwards from 1x1 box
headsize	head size
headangle	half angle of neck in terms of head
headloc	location of centre of head relative to origin with scale 1
necklength	neck length
shoulderwidth	shoulder width
shouldersize	size radius of arc for shoulder
armlength	arm length
armangle	angle of arm from horizontal
armwidth	width of arm
leglength	leg length
legangle	angle of leg from horizontal
legwidth	width of leg
torsolength	length of torso
...	other parameters passed to polygon()

Details

Draws a figure at a particular location. With defaults, has centre at origin and fits in 1x1 box.

Dimensions customisable

Value

invisibly returns co-ordinates

Examples

```
plot(0,xlim=c(-1,1),ylim=c(-1,1),type="n")
drawperson(0,0,1,col="yellow",border="red",lwd=3,lty=2)
```

---

drawprop	*drawprop*

Description

Illustrates a proportion as a set of people who are blue rather than red.

Usage

```
drawprop(prop, ci, nxy = 10, col1 = "maroon", col2 = "lightblue", ...)
```

Arguments

prop	the proportion to illustrate
ci	half the 95% CI width of the proportion.
nxy	illustrate on an n x n grid
col1	colour to put the 'in' proportion
col2	the other colour
...	passed to 'plot'

Details

Why anyone would want to think about a proportion this way is beyond the understanding of the authors, but the people have spoken.

Value

No return value, draws a figure

Examples

```
# See vignette
```

---

for_breakdown   *for_breakdown*

Description

For a given category (e.g., 'male', 'over 65') considers

1. all admissions for people in that category
2. all admissions for people in that category for which the SPARRA score was less than some threshold (e.g., false negatives

Usage

```
for_breakdown(
  decomp_table,
  group,
  threshold,
  inc_died = TRUE,
  ldiff = 0.005,
  ci = 0.95,
  xlimit = c(-0.05, 0.35),
  ylimit = c(-0.04, 0.04)
)
```

Arguments

`decomp_table`	matrix for group; see specification in description
`group`	name of group
`threshold`	cutoff, rounded to nearest 0.05
`inc_died`	set to TRUE to include a second panel showing 'death' type admissions
`ldiff`	specifically label points this far from xy line
`ci`	set to a value <1 to draw confidence intervals at that value, or FALSE to not draw confidence intervals.
`xlimit`	limits for x axis; default c(-0.05,0.35)
`ylimit`	limits for y axis; default c(-0.04,0.04)

Details

For each of these groups, we consider the breakdown of medical admission types. We then plot the frequency of admission types in group 1 against the difference in frequencies between group 1 and group 2 (group 2 minus group 1). An admission type which is relatively more common in group (1) indicates that, in the relevant category, the admission type tends to be associated with higher SPARRA scores (and is in a sense easier to predict). Such admission types will correspond to points below the line y=0. Admission types which are relatively more common in group 2 correspond to those which are relatively harder to predict. These correspond to points above the line y=0 Since points are close together, only those greater than a certain distance from 0 are marked.

Takes as an argument a matrix in which The matrix shows only data for the group in question Columns are named with the admission types to be plotted. Any admission types including the string 'Died' are counted as deaths If the matrix has N rows, these are interpreted as corresponding to N score quantiles in increasing order. The (i,j)th entry of the matrix is the number of people admitted for reason i with a score greater than or equal to (j-1)/N and less than (j/N) who are in that group

Value

ggplot figure (invisible)

Examples

```
# See vignette
```

---

getcal	*getcal()*

Description

Produces a set of points for a calibration plot.

Usage

```
getcal(
  y,
  ypred,
  n = 10,
  kernel = FALSE,
  kernel_sd = 0.05,
  alpha = 0.05,
  c0 = 0,
  c2 = 0.1
)
```

Arguments

y   class labels, 0/1 or logical

ypred   predictions Pr(Y=1), numeric vector

n   number of subintervals/points

kernel   set to TRUE to use kernel method

kernel_sd   kernel width for kernel method; see above

alpha   return a pointwise confidence envolope for conservative 1-alpha confidence interval

c0   for computing maximum bias; assume true covariance function is of the form $a_0 + a_1 x + a_2 x^2$, with $|a_0| < c_0$, $|a_2| < c_2$ (c1 does not matter)

c2   for computing maximum bias; assume true covariance function is of the form $a_0 + a_1 x + a_2 x^2$, with $|a_0| < c_0$, $|a_2| < c_2$ (c1 does not matter)

Details

Uses either a binning method or a kernel method to determine height of points.

In both methods, considers n equally spaced subintervals of (0,1)

Value

a list with components x (expected calibration), y (observed calibration), n (number of samples in bins, if relevant), lower/upper (confidence interval on y)

Examples

```
# See vignette
```

---

getprc   *getprc()*

Description

Comprehensive plotting function for precision-recall curve. Also calculates AUPRC and standard error.

Usage

```
getprc(y, ypred, cv = NULL, res = 100)
```

Arguments

y	class labels, 0/1 or logical
ypred	predictions Pr(Y=1), numeric vector
cv	cross-validation fold assignments, if relevant. Changes estimate of standard error.
res	resolution. Returns this many equally-spaced points along the curve. Set res to null to return all points.

Details

Rather than returning points corresponding to every cutoff, only returns a representative sample of equally-spaced points along the curve.

Does not plot anything. Object can be plotted in a default way.

Value

list containing: ppv, ppv for res points in every cv fold; sens, sensitivity for res points in every cv fold; auc, areas under the curve for each fold and average (note length is 1 greater than number of CV folds); se, standard error for AUC in each fold and standard error for average auc (note length is 1 greater than number of CV folds)

Examples

```
# See vignette
```

---

getroc	*getroc() Comprehensive plotting function for receiver-operator characteristic curve. Also calculates AUROC and standard error.*

Description

Rather than returning points corresponding to every cutoff, only returns a representative sample of equally-spaced points along the curve.

Usage

```
getroc(y, ypred, cv = NULL, res = 100)
```

Arguments

y	class labels, 0/1 or logical
ypred	predictions Pr(Y=1), numeric vector
cv	cross-validation fold assignments, if relevant. Changes estimate of standard error.
res	resolution. Returns this many equally-spaced points along the curve. Set res to null to return all points.

Details

SE of AUROC with no CV structure is from Hanley and McNeil 1982. SE of AUROC with CV
folds is from LeDell et al 2012

Does not plot anything. Object can be plotted in a default way.

Value

list containing: spec, specificity for res points in every cv fold; sens, sensitivity for res points in
every cv fold; auc, areas under the curve for each fold and average (note length is 1 greater than
number of CV folds); se, standard error for AUC in each fold and standard error for average auc
(note length is 1 greater than number of CV folds)

Examples

```
# See vignette
```

---

    groupmetric_2panel   *groupmetric_2panel Draws plots of a group fairness metric with a sec-*
                                      ond panel underneath

Description

groupmetric_2panel Draws plots of a group fairness metric with a second panel underneath

Usage

```
groupmetric_2panel(
  objs,
  labels = names(objs),
  col = 1:length(objs),
  yrange = NULL,
  ci_col = col,
  highlight = NULL,
  logscale = FALSE,
  lpos = c(1, 0),
  yrange_lower = NULL,
  legend_title = ""
)
```

Arguments

objs	list of fairness objects. Each should contain sub-objects 'x', 'y' and 'ci', which specify x and y values and half-widths of confidence intervals around y.
labels	labels to use in legend
col	line colours
yrange	limit of y axis; defaults to 0,1
ci_col	confidence envelope colours. These will be transparent.
highlight	if non-null, draw a point at a particular cutoff
logscale	if TRUE, draw with log-scale.

lpos	legend position; as accepted by ggplot legend.position
yrange_lower	y range for lower plot. If NULL, generates automatically
legend_title	title for legend, defaults to nothing

Value

Silently return ggplot object

Examples

```
# See vignette
```

---

group_fairness	*group_fairness*

Description

Estimates group fairness metric according to a specification vector of the form

Usage

```
group_fairness(
  specs,
  scores,
  target,
  group1,
  group2,
  cutoffs = seq(min(scores, na.rm = TRUE), max(scores, na.rm = TRUE), length = 100)
)
```

Arguments

specs	specification vector; see description
scores	vector of risk scores
target	vector of values of target (which risk score aims to predict)
group1	indices of group 1
group2	indices of group 2
cutoffs	score cutoffs at which to estimate metric (default 100 evenly-spaced)

Details

c(A1,B1,C1,A2,B2,C2)

encoding a probability

P(A1,B1,C1|A2,B2,C2)

where

A1/A2 are events coded by 1:'score>= cutoff'; 0: 'score<cutoff' and NA: 1/TRUE B1/B2 are events coded by 1:'target=TRUE'; 0: 'target=FALSE' and NA: 1/TRUE C1/C2 are events coded by 1:'group=g'; and NA: 1/TRUE

For example, specs=c(NA,1,NA,0,NA,1) would encode false omission rate:

P(target=TRUE|score<cutoff,group=g)

Value

matrix of dimension length(cutoffs)x4, with (i,2g-1)th entry the relevant fairness metric for group g at the ith cutoff value and (i,2g)th entry the approximate standard error of the (i,2g-1)th entry

Examples

```
# See vignette
```

---

integral	*integral() Quick form for trapezoidal integration over range of x*

Description

integral() Quick form for trapezoidal integration over range of x

Usage

```
integral(x, y = NULL)
```

Arguments

x	x co-ordinates, or nx2 matrix of points
y	y co-ordinates

Value

trapezoidal estimate of integral of the xth value of y over range of x.

---

logistic	*Logistic*

Description

Logistic function: 1/(1+exp(-x))

Usage

```
logistic(x)
```

Arguments

x	argument

Value

value of logistic(x)

Examples

```
# Plot
x=seq(-5,5,length=1000)
plot(x,logistic(x),type="l")
```

---

logit   *Logit*

Description

Logit function: -log((1/x)-1)

Usage

```
logit(x)
```

Arguments

x   argument

Value

value of logit(x); na if x is outside (0,1)

Examples

```
# Plot
x=seq(0,1,length=100)
plot(x,logit(x),type="l")

# Logit and logistic are inverses
x=seq(-5,5,length=1000)
plot(x,logit(logistic(x)),type="l")
```

---

phs_colours   *phs_colours*

Description

Copied from github, "Public-Health-Scotland/phsstyles". Public Health Scotland colour scheme. Internal function.

Usage

```
phs_colours(colourname = NULL, keep_names = FALSE)
```

Arguments

colourname   name of colour; usually something like phs-blue. If NULL returns all colours.

keep_names   keep names of colours in return list. Defaults to false.

Value

vector of colours, optionally with names.

---

plot.sparraCAL   *Plot function for class sparraCAL*

Description

Plot function for class sparraCAL

Usage

```
## S3 method for class 'sparraCAL'
plot(
  x,
  cols = rep(phs_colours("phs-blue"), dim(x$sens)[1]),
  add = FALSE,
  add_xy_line = TRUE,
  ...
)
```

Arguments

x	output from getcal()
cols	colour to draw lines
add	set to FALSE to add to existing plot
add_xy_line	set to TRUE to draw an X-Y reference line.
...	passed to lines()

Value

No return value, draws a figure

Examples

```
# See vignette
```

---

plot.sparraPRC   *Plot function for class above*

Description

Plot function for class above

Usage

```
## S3 method for class 'sparraPRC'
plot(
  x,
  addauc = FALSE,
  cols = rep(phs_colours("phs-blue"), dim(x$sens)[1]),
  ...
)
```

Arguments

x	output from getprc()
addauc	set to TRUE to add text to the plot showing the (mean) AUC and SE.
cols	colour to draw lines
...	passed to plot()

Value

No return value, draws a figure

Examples

```
# See vignette
```

---

plot.sparraROC *Plot function for class sparraROC*

Description

Plot function for class sparraROC

Usage

```
## S3 method for class 'sparraROC'
plot(
  x,
  addauc = FALSE,
  cols = rep(phs_colours("phs-blue"), dim(x$sens)[1]),
  ...
)
```

Arguments

x	output from getroc()
addauc	set to TRUE to add text to the plot showing the (mean) AUC and SE.
cols	colour to draw lines
...	passed to plot()

Value

No return value, draws a figure

Examples

```
# See vignette
```

---

plot_decomp   *plot_decomp*

Description

Plots a bar graph of decomposition of FORP by cause of admission

Usage

```
plot_decomp(decomp1, decomp2, threshold, labels, inc_died = TRUE)
```

Arguments

decomp1   matrix for first group; see specification in description

decomp2   matrix for second group; see specification in description

threshold   score threshold to plot (between 0 and 1)

labels   labels for group 1 and group 2

inc_died   set to TRUE to include a second panel showing 'death' type admissions

Details

Takes two matrices as input with the following specifications: Each matrix corresponds to one group Columns are named with the admission types to be plotted. Any admission types including the string 'Died' are counted as deaths If the matrix has N rows, these are interpreted as corresponding to N score quantiles in increasing order. The (i,j)th entry of the matrix is the number of people admitted for reason i with a score greater than or equal to (j-1)/N and less than (j/N) who are in that group

Value

Silently return ggplot object

Examples

```
# See vignette
```

---

prc_2panel   *prc_2panel Draws a PRC curve (with legend) with a second panel*
                              underneath showing precision difference.

Description

prc_2panel Draws a PRC curve (with legend) with a second panel underneath showing precision difference.

Usage

```
prc_2panel(
  prcs,
  labels = names(prcs),
  col = 1:length(prcs),
  highlight = NULL,
  yrange_lower = NULL,
  legend_title = ""
)
```

Arguments

prcs	list of sparraPRC objects.
labels	labels to use in legend
col	line colours
highlight	if non-null, draw a point at a particular cutoff
yrange_lower	y range for lower plot. If NULL, generates automatically
legend_title	title for legend, defaults to nothing

Value

Silently return ggplot object

Examples

```
# See vignette
```

---

roc_2panel	*roc_2panel Draws a ROC curve (with legend) with a second panel underneath showing sensitivity difference.*

Description

roc_2panel Draws a ROC curve (with legend) with a second panel underneath showing sensitivity difference.

Usage

```
roc_2panel(
  rocs,
  labels = names(rocs),
  col = 1:length(rocs),
  xy_col = phs_colours("phs-magenta"),
  highlight = NULL,
  yrange_lower = NULL,
  legend_title = ""
)
```

Arguments

`rocs`	list of sparraROC objects
`labels`	labels to use in legend; default to names of rocs.
`col`	line colours
`xy_col`	line colour for x-y line, defaults to red
`highlight`	if non-null, add a point at this cutoff
`yrange_lower`	y range for lower plot. If NULL, generates automatically
`legend_title`	title for legend, defaults to nothing

Value

Invisibly returns plot as ggplot object

Examples

```
# See vignette
```

---

sim_pop_data	*sim_pop_data*

Description

Simulates population data with a reasonably realistic joint distribution

Usage

```
sim_pop_data(
  npop,
  coef_adjust = 4,
  offset = 1,
  vcor = NULL,
  coefs = c(2, 1, 0, 5, 3, 0, 0),
  seed = 12345,
  incl_id = TRUE,
  incl_reason = TRUE
)
```

Arguments

`npop`	population size
`coef_adjust`	inverse scale for all (true) coefficients (default 4): lower means that hospital admissions are more predictable from covariates.
`offset`	offset for logistic model (default 1): higher means a lower overall prevalence of admission
`vcor`	a valid 5x5 correlation matrix (default NULL), giving correlation between variables. If 'NULL', values roughly represents realistic data.
`coefs`	coefficients of age, male sex, non-white ethnicity, number of previous admissions, and deprivation decile on hospital admissions, Default (2,1,0,5,3). Divided through by coef_adjust.

seed	random seed (default 12345)
incl_id	include an ID column (default TRUE)
incl_reason	include a column indicating reason for admission.

Details

Simulates data for a range of people for the variables

- Age (`age`)
- Sex (`sexM`; 1 if male)
- Race/ethnicity (`raceNW`: 1 if non-white ethnicity)
- Number of previous hospital admissions (`PrevAdm`)
- Deprivation decile (`SIMD`: 1 most deprived, 10 least deprived. NOTE - opposite to English IMD)
- Urban-rural residence status (`urban_rural`: 1 for rural)
- Mainland-island residence status (`mainland_island`: 1 for island)
- Hospital admission (`target`: 1/TRUE if admitted to hospital in year following prediction date)

Can optionally add an ID column.

Optionally includes an admission reason for samples with `target=1`. These admission reasons roughly correspond to the first letters of ICD10 categories, and can either correspond to an admission or death. Admission reasons are simulated with a non-constant multinomial distribution which varies across age/sex/ethnicity/urban-rural/mainland-island/PrevAdm values in a randomly- chosen way. The distributions of admission reasons are *not* however chosen to reflect real distributions, nor are systematic changes in commonality of admission types across categories intended to appear realistic.

Value

data frame with realistic values.

Examples

```
# Simulate data
dat=sim_pop_data(10000)
cor(dat[,1:7])

# See vignette
```

Index